



WheelsEye

Wheelseye Online Truck Booking Machine Learning Project

Project By : PRASAD JADHAV

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: df = pd.read_csv('wheelseye_data.csv')
pd.set_option('display.max_columns',20)
print(df.shape)
```

(10000, 17)

```
In [4]: df.head()
```

```
Out[4]:
```

	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight	tru
0	ORDI000	CUS9726	Chennai	Delhi	1144	Open Truck	1.61 Tons	
1	ORDI001	CUS7464	Chennai	Pune	1912	Open Truck	1.98 Tons	
2	ORDI002	CUS4668	Kanpur	Indore	1145	Open Truck	7.28 Tons	
3	ORDI003	CUS2876	Indore	Chennai	1247	Open Truck	8.69 Tons	
4	ORDI004	CUS4137	Indore	Kanpur	569	Small Pickup Truck	7.6 Tons	

```
In [6]: df.tail()
```

Out[6]:	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight
9995	ORD10995	CUS9223	Chennai	Ludhiana	643	Container	9.29 Ton
9996	ORD10996	CUS6562	Ludhiana	Pune	917	Small Pickup Truck	2.71 Ton
9997	ORD10997	CUS3532	Pune	Delhi	1687	Container	9.76 Ton
9998	ORD10998	CUS3105	Chennai	Indore	484	Open Truck	3.93 Ton
9999	ORD10999	CUS2121	Pune	Ahmedabad	1018	Small Pickup Truck	5.72 Ton

Dataset Overview

- `order_id`: Unique identifier for each transport order (e.g., ORD12345)
- `customer_id`: Unique identifier for each customer (e.g., CUS5678)
- `city_from`: Starting city (e.g., Delhi, Mumbai)
- `city_to`: Destination city (e.g., Kanpur, Chennai)
- `route_distance_km`: The total distance in kilometers between the start and destination city
- `vehicle_type`: Type of vehicle used for transport (Open Truck, Container, Small Pickup Truck)
- `material_weight`: Weight of the material being transported (in Tons or Kg)
- `truck_length_ft`: Length of the truck (7, 8, 14, 17, 19, 20, 22 ft)
- `truck_height_ft`: Height of the truck (6.0 ft for smaller trucks, 8.0 ft for larger trucks)
- `business_category`: Category of the customer (Manufacturer, Transporter, Truck Owner, Individual, Other)
- `weather_condition`: Weather conditions during the transport (e.g., Sunny, Rainy, Cloudy)
- `traffic_condition`: Traffic conditions on the route (e.g., Light, Moderate, Heavy)
- `driver_rating`: Rating of the driver out of 5 stars (e.g., 4.5)
- `customer_rating`: Customer's rating of the service out of 5 stars (e.g., 4.8)
- `order_price`: Price of the transport service (to be fetched based on route, weight, vehicle type, etc.)
- `delivery_time_hours`: The time taken for the transport service in hours
- `is_delayed`: Boolean indicating if the delivery was delayed (True or False)

Data Exploration

In [7]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   order_id              10000 non-null  object
1   customer_id           10000 non-null  object
2   city_from             10000 non-null  object
3   city_to               10000 non-null  object
4   route_distance_km     10000 non-null  int64
5   vehicle_type          10000 non-null  object
6   material_weight       10000 non-null  object
7   truck_length_ft       10000 non-null  int64
8   truck_height_ft       10000 non-null  float64
9   business_category     10000 non-null  object
10  weather_condition     10000 non-null  object
11  traffic_condition     10000 non-null  object
12  driver_rating         10000 non-null  float64
13  customer_rating       10000 non-null  float64
14  order_price           10000 non-null  object
15  delivery_time_hours   10000 non-null  float64
16  is_delayed            10000 non-null  bool
dtypes: bool(1), float64(4), int64(2), object(10)
memory usage: 1.2+ MB

```

```
In [8]: df.isnull().sum()
```

```

Out[8]: order_id              0
        customer_id          0
        city_from            0
        city_to              0
        route_distance_km    0
        vehicle_type         0
        material_weight      0
        truck_length_ft      0
        truck_height_ft      0
        business_category    0
        weather_condition    0
        traffic_condition    0
        driver_rating        0
        customer_rating      0
        order_price          0
        delivery_time_hours  0
        is_delayed           0
dtype: int64

```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: np.int64(0)
```

```
In [10]: df.describe()
```

```
Out[10]:
```

	route_distance_km	truck_length_ft	truck_height_ft	driver_rating	customer_rating	delivery-
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	1048.489900	15.33730	6.86220	4.003230	3.998540	
std	547.719651	5.45189	0.99051	0.576817	0.580458	
min	100.000000	7.00000	6.00000	3.000000	3.000000	
25%	576.000000	8.00000	6.00000	3.500000	3.500000	
50%	1048.500000	17.00000	6.00000	4.000000	4.000000	
75%	1526.000000	20.00000	8.00000	4.500000	4.500000	
max	2000.000000	22.00000	8.00000	5.000000	5.000000	

```
In [11]: df.head()
```

```
Out[11]:
```

	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight	tru
0	ORDI000	CVS9726	Chennai	Delhi	1144	Open Truck	1.61 Tons	
1	ORDI001	CVS7464	Chennai	Pune	1912	Open Truck	1.98 Tons	
2	ORDI002	CVS4668	Kanpur	Indore	1145	Open Truck	7.28 Tons	
3	ORDI003	CVS2876	Indore	Chennai	1247	Open Truck	8.69 Tons	
4	ORDI004	CVS4137	Indore	Kanpur	569	Small Pickup Truck	7.6 Tons	

Step by setp columns price predict ml problem # city_from city_to, Vehicle Type, Material Weight, Truck Length, Truck Height, Business Category:,

```
In [12]: num_cols = [x for x in df.columns if df[x].dtypes != 'float64']

for col in num_cols:
    print(f"Value counts for column '{col}':")
    print(df[col].value_counts())
    print("\n" + "_"*40 + "\n")
```

```
Value counts for column 'order_id':
order_id
ORD1000      1
ORD7670      1
ORD7663      1
ORD7664      1
ORD7665      1
..
ORD4333      1
ORD4334      1
ORD4335      1
ORD4336      1
ORD10999     1
Name: count, Length: 10000, dtype: int64
```

```
Value counts for column 'customer_id':
customer_id
CUS1076      6
CUS8970      6
CUS3108      6
CUS9981      6
CUS4917      6
..
CUS5144      1
CUS1596      1
CUS7660      1
CUS7743      1
CUS6562      1
Name: count, Length: 6033, dtype: int64
```

```
Value counts for column 'city_from':
city_from
Kanpur      1153
Indore      1136
Mumbai      1133
Ludhiana    1121
Pune        1116
Ahmedabad   1112
Chennai     1100
Delhi       1089
Kolkata     1040
Name: count, dtype: int64
```

```
Value counts for column 'city_to':
city_to
Chennai     1181
Kanpur      1148
Ludhiana    1140
Pune        1109
Ahmedabad   1108
```

```
Indore      1096
Delhi       1084
Kolkata     1067
Mumbai      1067
Name: count, dtype: int64
```

```
Value counts for column 'route_distance_km':
route_distance_km
1436      19
1545      14
727       13
1804      13
675       13
..
667       1
368       1
890       1
748       1
1627      1
Name: count, Length: 1888, dtype: int64
```

```
Value counts for column 'vehicle_type':
vehicle_type
Small Pickup Truck    3377
Container              3329
Open Truck            3294
Name: count, dtype: int64
```

```
Value counts for column 'material_weight':
material_weight
0.64 Tons      22
6.57 Tons      21
7.6 Tons       21
1.44 Tons      21
9.21 Tons      19
..
0.5 Tons       3
3.05 Tons      3
3.57 Tons      3
5.77 Tons      3
4.13 Tons      1
Name: count, Length: 951, dtype: int64
```

```
Value counts for column 'truck_length_ft':
truck_length_ft
22      1512
17      1444
8       1435
```

```
19    1432
14    1409
7      1401
20    1367
Name: count, dtype: int64
```

```
Value counts for column 'business_category':
business_category
Other          2066
Transporter    2045
Truck Owner    2038
Individual     1945
Manufacturer   1906
Name: count, dtype: int64
```

```
Value counts for column 'weather_condition':
weather_condition
Sunny      3371
Rainy      3360
Cloudy     3269
Name: count, dtype: int64
```

```
Value counts for column 'traffic_condition':
traffic_condition
Heavy      3381
Moderate   3321
Light      3298
Name: count, dtype: int64
```

```
Value counts for column 'order_price':
order_price
₹21979.46    2
₹23487.09    2
₹21104.62    2
₹13382.77    2
₹29033.66    2
..
₹25786.97    1
₹33208.61    1
₹26283.9     1
₹21896.34    1
₹18041.51    1
Name: count, Length: 9977, dtype: int64
```

```
Value counts for column 'is_delayed':
is_delayed
```

```
False    5052  
True      4948  
Name: count, dtype: int64
```

```
In [13]: cat_cols = [col for col in df.columns if df[col].dtype == 'object' or df[col]  
  
for col in cat_cols:  
    print(f"Value counts for column '{col}':")  
    print(df[col].value_counts())  
    print("\n" + "_"*40 + "\n")
```



```
Value counts for column 'order_id':
order_id
ORD1000      1
ORD7670      1
ORD7663      1
ORD7664      1
ORD7665      1
..
ORD4333      1
ORD4334      1
ORD4335      1
ORD4336      1
ORD10999     1
Name: count, Length: 10000, dtype: int64
```

```
Value counts for column 'customer_id':
customer_id
CUS1076      6
CUS8970      6
CUS3108      6
CUS9981      6
CUS4917      6
..
CUS5144      1
CUS1596      1
CUS7660      1
CUS7743      1
CUS6562      1
Name: count, Length: 6033, dtype: int64
```

```
Value counts for column 'city_from':
city_from
Kanpur      1153
Indore      1136
Mumbai      1133
Ludhiana    1121
Pune        1116
Ahmedabad   1112
Chennai     1100
Delhi       1089
Kolkata     1040
Name: count, dtype: int64
```

```
Value counts for column 'city_to':
city_to
Chennai     1181
Kanpur      1148
Ludhiana    1140
Pune        1109
Ahmedabad   1108
```

```
Indore      1096
Delhi       1084
Kolkata     1067
Mumbai      1067
Name: count, dtype: int64
```

```
Value counts for column 'vehicle_type':
vehicle_type
Small Pickup Truck    3377
Container              3329
Open Truck            3294
Name: count, dtype: int64
```

```
Value counts for column 'material_weight':
material_weight
0.64 Tons    22
6.57 Tons    21
7.6 Tons     21
1.44 Tons    21
9.21 Tons    19
..
0.5 Tons     3
3.05 Tons     3
3.57 Tons     3
5.77 Tons     3
4.13 Tons     1
Name: count, Length: 951, dtype: int64
```

```
Value counts for column 'business_category':
business_category
Other          2066
Transporter     2045
Truck Owner     2038
Individual      1945
Manufacturer    1906
Name: count, dtype: int64
```

```
Value counts for column 'weather_condition':
weather_condition
Sunny          3371
Rainy          3360
Cloudy         3269
Name: count, dtype: int64
```

```
Value counts for column 'traffic_condition':
traffic_condition
```


```
Heavy      3381
Moderate   3321
Light      3298
Name: count, dtype: int64
```

```
Value counts for column 'order_price':
order_price
₹21979.46    2
₹23487.09    2
₹21104.62    2
₹13382.77    2
₹29033.66    2
..
₹25786.97    1
₹33208.61    1
₹26283.9     1
₹21896.34    1
₹18041.51    1
Name: count, Length: 9977, dtype: int64
```

```
In [14]: num_features = df.select_dtypes(include = ['int64', 'float64']).dtypes.index
```

```
In [15]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=15)

for i in range(0,len(num_features)):
    plt.subplot(10,4,i+1)
    sns.kdeplot(x=df[num_features[i]],shade=True,color='brown')
    plt.tight_layout()
```



```
In [ ]: plt.figure(figsize = (15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=20)

for i in range(0,len(num_features)):
    plt.subplot(10,5,i+1)
    sns.boxplot(data=df,x=num_features[i],color='brown')
    plt.xlabel(num_features[i])
    plt.tight_layout()
```

```
In [ ]: cat_features = df.select_dtypes(include='object').dtypes.index
```

```
In [ ]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of Features',fontweight='bold',fontsize=15

for i in range(0,len(cat_features)):
    plt.subplot(10,4,i+1)
    sns.countplot(x=df[cat_features[i]],color='brown')
    plt.tight_layout()
```

```
In [ ]: # Example Machine Learning Use Cases:
```

```
In [ ]: # Preprocessing: Convert categorical variables into dummy/one-hot encoding
# df = pd.get_dummies(df, columns=['city_from', 'city_to', 'vehicle_type', 'business_category'])
```

```
In [16]: # Define mappings for each categorical feature
city_mapping = {city: idx for idx, city in enumerate(df['city_from'].unique())}
vehicle_mapping = {vehicle: idx for idx, vehicle in enumerate(df['vehicle_type'].unique())}
business_mapping = {business: idx for idx, business in enumerate(df['business_category'].unique())}

# Apply the mappings to the columns
df['city_from'] = df['city_from'].map(city_mapping)
df['city_to'] = df['city_to'].map(city_mapping) # Assuming city_to has similar unique values
df['vehicle_type'] = df['vehicle_type'].map(vehicle_mapping)
df['business_category'] = df['business_category'].map(business_mapping)
```

```
In [17]: # Remove the word "Tons" from the material_weight column and convert to float
df['material_weight'] = df['material_weight'].str.replace("Tons", "").astype(float)
```

```
In [18]: # Define mappings for the categorical features
weather_mapping = {condition: idx for idx, condition in enumerate(df['weather_condition'].unique())}
traffic_mapping = {condition: idx for idx, condition in enumerate(df['traffic_condition'].unique())}
is_delayed_mapping = {'True': 1, 'False': 0} # Assuming is_delayed has 'True' and 'False' values

# Apply the mappings to each column
df['weather_condition'] = df['weather_condition'].map(weather_mapping)
df['traffic_condition'] = df['traffic_condition'].map(traffic_mapping)
df['is_delayed'] = df['is_delayed'].map(is_delayed_mapping)
```

```
In [19]: # df['is_delayed'] = df['is_delayed'].map({'False':0, 'True':1}).astype(int)
```

```
In [20]: df['is_delayed'] = df['is_delayed'].astype(int)
```

Price Prediction (Regression)

```
In [22]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
```

```
In [23]: price_features = df[['city_from', 'city_to', 'vehicle_type', 'material_weight', 'business_category', 'weather_condition', 'traffic_condition', 'is_delayed']]
```

```
In [24]: # Feature Engineering: Remove unnecessary columns and extract target (price)
df['order_price'] = df['order_price'].replace(['₹,'], '', regex=True).astype(float)
```

```
# X = df.drop(['order_id', 'customer_id', 'order_price', 'is_delayed'], axis=1)
X = price_features
y = df['order_price'] # Target (Price)
```

```
In [25]: # Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [26]: # Model Training: Random Forest Regressor
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
Out[26]: RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
In [27]: # Prediction and Evaluation
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Price Prediction RMSE: {rmse}")
```

Price Prediction RMSE: 8673.3491955699

Delivery Time Estimation (Regression)

```
In [28]: from sklearn.tree import DecisionTreeRegressor
```

```
In [29]: # Preprocessing: Use the same df but now we focus on predicting delivery time
X = df.drop(['order_id', 'customer_id', 'delivery_time_hours', 'is_delayed'], axis=1)
y = df['delivery_time_hours'] # Target (Delivery Time)
```

```
In [30]: # Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [31]: # Model Training: Random Forest Regressor for Delivery Time
model = DecisionTreeRegressor(random_state=42)
model.fit(X_train, y_train)
```

```
Out[31]: DecisionTreeRegressor
DecisionTreeRegressor(random_state=42)
```

```
In [32]: # Prediction and Evaluation
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Delivery Time Estimation RMSE: {rmse}")
```

Delivery Time Estimation RMSE: 5.053821865677499

Customer Satisfaction Analysis
(Regression/Classification)

```
In [33]: from sklearn.linear_model import LinearRegression
```

```
In [34]: # Regression (Predict Customer Rating)
```

```
In [35]: # Preprocessing: Use the same df but now focus on predicting customer rating
X = df.drop(['order_id', 'customer_id', 'customer_rating', 'is_delayed'], ax
y = df['customer_rating'] # Target (Customer Rating)
```

```
In [36]: # Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [37]: # Model Training: Random Forest Regressor for Customer Rating
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[37]: ▾ LinearRegression ⓘ ?
LinearRegression()
```

```
In [38]: # Prediction and Evaluation
y_pred = model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Customer Rating Prediction RMSE: {rmse}")
```

Customer Rating Prediction RMSE: 0.5816678982940906

```
In [39]: # Classification (High/Low Satisfaction)
```

```
In [40]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [41]: # Create a binary column based on customer rating threshold
df['satisfied'] = df['customer_rating'] >= 4.0 # True for ratings >= 4, Fal
```

```
In [42]: # Preprocessing: Drop unnecessary columns
X = df.drop(['order_id', 'customer_id', 'customer_rating', 'is_delayed', 'sa
y = df['satisfied'].astype(int) # Target (Satisfied/Not Satisfied)
```

```
In [43]: # Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [44]: model = LogisticRegression()
model.fit(X_train, y_train)
```

```
Out[44]: ▾ LogisticRegression ⓘ ?
LogisticRegression()
```

```
In [45]: # Prediction and Evaluation
y_pred = model.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Customer Satisfaction Classification Accuracy: {accuracy}")
```

Customer Satisfaction Classification Accuracy: 0.5155

Delay Prediction (Classification)

```
In [46]: from sklearn.ensemble import RandomForestClassifier
```

```
In [47]: # Preprocessing: Predict delay status (is_delayed) as binary classification
X = df.drop(['order_id', 'customer_id', 'is_delayed'], axis=1) # Features
y = df['is_delayed'].astype(int) # Target (0: On-time, 1: Delayed)
```

```
In [48]: # Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```
In [49]: # Model Training: Random Forest Classifier for Delay Prediction
model = RandomForestClassifier(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
Out[49]: ▼      RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
In [50]: # Prediction and Evaluation
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Delay Prediction Accuracy: {accuracy}")
```

Delay Prediction Accuracy: 0.504

```
In [51]: df.head()
```

```
Out[51]:
```

	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight	trun
0	ORDI000	CVS9726	0	3	1144	0	1.61	
1	ORDI001	CVS7464	0	4	1912	0	1.98	
2	ORDI002	CVS4668	1	2	1145	0	7.28	
3	ORDI003	CVS2876	2	0	1247	0	8.69	
4	ORDI004	CVS4137	2	1	569	1	7.60	

```
In [52]: from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error

X = price_features # 'price_features' should be a DataFrame with relevant p
y = df['order_price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
```

```

# Train an XGBRegressor with default hyperparameters
xg_model = XGBRegressor()
xg_model.fit(X_train, y_train)

# Prediction and Evaluation
y_pred = xg_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"Price Prediction RMSE: {rmse}")

```

Price Prediction RMSE: 8963.408836849994

```

In [53]: from xgboost import XGBRegressor
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

# Ensure price_features and df['order_price'] are defined
# Example: df['order_price'] = df['order_price'].replace('[₹,]', '', regex=True)
# X = price_features

# Define features and target variable
X = price_features # 'price_features' should be a DataFrame with relevant p
y = df['order_price']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Train an XGBRegressor with default hyperparameters
model = XGBRegressor()
model.fit(X_train, y_train)

# Evaluate the model on the training set
train_preds = model.predict(X_train)
train_mse = mean_squared_error(y_train, train_preds)
train_rmse = np.sqrt(train_mse)

# Evaluate the model on the testing set
test_preds = model.predict(X_test)
test_mse = mean_squared_error(y_test, test_preds)
test_rmse = np.sqrt(test_mse)

print(f"Training RMSE: {train_rmse:.2f}")
print(f"Testing RMSE: {test_rmse:.2f}")

```

Training RMSE: 5779.81

Testing RMSE: 8963.41

```

In [54]: # XGBoost for Price Prediction (Regression)

import xgboost as xgb
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split

# Define features and target variable
X = price_features # price_features should contain only the predictor varia
y = df['order_price']

```



```

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Instantiate XGBoost Regressor
xgb_reg = xgb.XGBRegressor(objective='reg:squarederror', random_state=42)

# Define a comprehensive hyperparameter grid for fine-tuning
param_grid = {
    'n_estimators': [100, 300, 500, 700],
    'max_depth': [3, 5, 7, 9],
    'learning_rate': [0.01, 0.05, 0.1, 0.2, 0.3],
    'subsample': [0.6, 0.7, 0.8, 0.9, 1.0],
    'colsample_bytree': [0.5, 0.7, 0.8, 0.9, 1.0],
    'gamma': [0, 0.1, 0.2, 0.3, 0.4],
    'min_child_weight': [1, 3, 5, 7]
}

# Set up RandomizedSearchCV with cross-validation
random_search = RandomizedSearchCV(
    estimator=xgb_reg,
    param_distributions=param_grid,
    n_iter=20, # Increased for more comprehensive search
    scoring='neg_root_mean_squared_error',
    n_jobs=-1,
    cv=3,
    verbose=2,
    random_state=42
)

# Fit the model to find the best hyperparameters
random_search.fit(X_train, y_train)

# Display the best parameters found
print(f"Best parameters: {random_search.best_params_}")

# Predict on the test set and evaluate performance
y_pred = random_search.best_estimator_.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"XGBoost Price Prediction RMSE: {rmse}")

```

Fitting 3 folds for each of 20 candidates, totalling 60 fits

Best parameters: {'subsample': 0.9, 'n_estimators': 100, 'min_child_weight': 3, 'max_depth': 3, 'learning_rate': 0.01, 'gamma': 0, 'colsample_bytree': 0.8}

XGBoost Price Prediction RMSE: 8336.01500656909

```

In [55]: # XGBoost for Delay Prediction (Classification)

import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score

# Preprocessing for delay classification
X = df.drop(['order_id', 'customer_id', 'is_delayed'], axis=1)
y = df['is_delayed'].astype(int)

```

```

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Hyperparameter tuning using RandomizedSearchCV with XGBoost
xgb_clf = xgb.XGBClassifier()

# Define the hyperparameters and the ranges for tuning
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0]
}

# RandomizedSearchCV
random_search = RandomizedSearchCV(xgb_clf, param_distributions=param_grid,
                                   n_jobs=-1, cv=3, verbose=2, random_state=

# Fit the model
random_search.fit(X_train, y_train)

# Best parameters
print(f"Best parameters: {random_search.best_params_}")

# Evaluate the model on test data
y_pred = random_search.best_estimator_.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"XGBoost Delay Prediction Accuracy: {accuracy}")

```

Fitting 3 folds for each of 10 candidates, totalling 30 fits
 Best parameters: {'subsample': 1.0, 'n_estimators': 300, 'max_depth': 7, 'learning_rate': 0.1, 'colsample_bytree': 1.0}
 XGBoost Delay Prediction Accuracy: 0.5085

```

In [56]: # LightGBM for Price Prediction (Regression)

import lightgbm as lgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import mean_squared_error

# Preprocessing as before
X = price_features # price_features should contain only the predictor variables
y = df['order_price']

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Hyperparameter tuning using RandomizedSearchCV with LightGBM
lgb_reg = lgb.LGBMRegressor()

# Define the hyperparameters and the ranges for tuning
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7],

```

```

    'learning_rate': [0.01, 0.1, 0.2],
    'num_leaves': [31, 50, 100],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0]
}

# RandomizedSearchCV
random_search = RandomizedSearchCV(lgb_reg, param_distributions=param_grid,
                                   n_jobs=-1, cv=3, verbose=2, random_state=

# Fit the model
random_search.fit(X_train, y_train)

# Best parameters
print(f"Best parameters: {random_search.best_params_}")

# Evaluate the model on test data
y_pred = random_search.best_estimator_.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"LightGBM Price Prediction RMSE: {rmse}")

```

[illegible]


```

cat_reg = CatBoostRegressor(verbose=0)

# Define the hyperparameters and the ranges for tuning
param_grid = {
    'iterations': [100, 200, 300],
    'depth': [4, 6, 8],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'l2_leaf_reg': [3, 5, 7],
    'bagging_temperature': [0, 1, 2], # Specific for regression to control
    'random_strength': [1, 2, 5]      # Controls the amount of randomness i
}

# RandomizedSearchCV
random_search = RandomizedSearchCV(cat_reg, param_distributions=param_grid,
                                   n_jobs=-1, cv=3, verbose=2, random_state=

# Fit the model
random_search.fit(X_train, y_train)

# Best parameters
print(f"Best parameters: {random_search.best_params_}")

# Evaluate the model on test data
y_pred = random_search.best_estimator_.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f"CatBoost Delay Prediction RMSE: {rmse}")

```

```

In [ ]: # CatBoost for Delay Prediction (Classification)

from catboost import CatBoostClassifier
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score

# Preprocessing for CatBoost (no need to one-hot encode)
X = df.drop(['order_id', 'customer_id', 'is_delayed'], axis=1)
y = df['is_delayed'].astype(int)

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Hyperparameter tuning using RandomizedSearchCV with CatBoost
cat_clf = CatBoostClassifier(verbose=0)

# Define the hyperparameters and the ranges for tuning
param_grid = {
    'iterations': [100, 200, 300],
    'depth': [4, 6, 8],
    'learning_rate': [0.01, 0.1, 0.2],
    'l2_leaf_reg': [3, 5, 7]
}

# RandomizedSearchCV
random_search = RandomizedSearchCV(cat_clf, param_distributions=param_grid,
                                   n_jobs=-1, cv=3, verbose=2, random_state=

# Fit the model

```

```
random_search.fit(X_train, y_train)

# Best parameters
print(f"Best parameters: {random_search.best_params_}")

# Evaluate the model on test data
y_pred = random_search.best_estimator_.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"CatBoost Delay Prediction Accuracy: {accuracy}")
```

```
In [59]: # Cross-validation

from sklearn.model_selection import cross_val_score

# Use cross-validation to evaluate the model
xgb_best = random_search.best_estimator_
scores = cross_val_score(xgb_best, X_train, y_train, cv=5, scoring='neg_root')
print(f"Cross-Validation RMSE: {-scores.mean():.4f}")
```

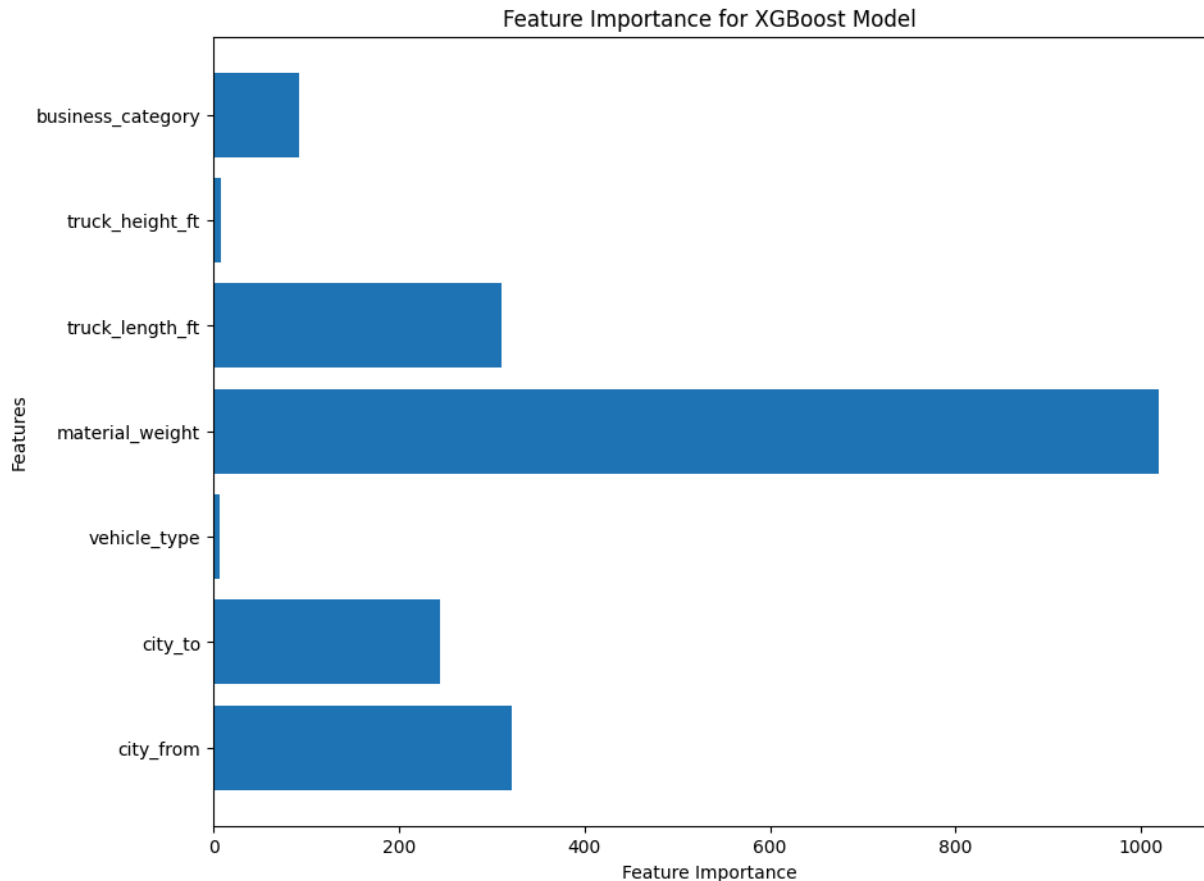

[illegible]

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf  
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

[illegible]

[illegible]


```
plt.title('Feature Importance for XGBoost Model')
plt.show()
```



```
In [62]: # Ensemble Methods

from sklearn.ensemble import StackingRegressor
from sklearn.linear_model import LinearRegression

# Define base models
estimators = [
    ('xgb', xgb.XGBRegressor(**random_search.best_params_)),
    ('lgb', lgb.LGBMRegressor(**random_search.best_params_)),
    # ('cat', CatBoostRegressor(verbose=0))
]

# Stacking regressor with Linear Regression as the final estimator
stacking_reg = StackingRegressor(estimators=estimators, final_estimator=Line

# Fit the stacking model
stacking_reg.fit(X_train, y_train)

# Predict and evaluate
y_pred_stack = stacking_reg.predict(X_test)
stack_rmse = np.sqrt(mean_squared_error(y_test, y_pred_stack))
print(f"Stacking RMSE: {stack_rmse:.4f}")
```


[illegible]

[illegible]

[illegible]

[illegible]

[illegible]


```
final_rmse = np.sqrt(mean_squared_error(y_test, y_final_pred))  
print(f"Final Model RMSE: {final_rmse:.4f}")
```



```

import joblib
import pandas as pd

app = Flask(__name__)

# Load the model
model = joblib.load('final_model_xgboost.pkl')

@app.route('/predict', methods=['POST'])
def predict():
    data = request.get_json() # Get data from request
    df = pd.DataFrame(data) # Convert to DataFrame
    prediction = model.predict(df) # Make prediction
    return jsonify({'prediction': list(prediction)})

if __name__ == '__main__':
    app.run(debug=True)

```

* Serving Flask app '__main__'

* Debug mode: on

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

* Running on http://127.0.0.1:5000

Press CTRL+C to quit

* Restarting with watchdog (windowsapi)

An exception has occurred, use %tb to see the full traceback.

SystemExit: 1

In [70]: *# Flask Deployment*

```

from flask import Flask, request, jsonify
import joblib
import pandas as pd

app = Flask(__name__)

# Load the trained model
try:
    model = joblib.load('final_model_xgboost.pkl')
except Exception as e:
    print(f"Error loading model: {str(e)}")

@app.route('/predict', methods=['POST'])
def predict():
    try:
        # Get the JSON data from the POST request
        data = request.get_json()

        # Convert the JSON data into a Pandas DataFrame
        df = pd.DataFrame([data])

        # Ensure the data structure is correct before making predictions
        if df.empty:
            return jsonify({"error": "Empty input data"})

```

```

        # Make predictions using the model
        prediction = model.predict(df)

        # Return the prediction as a JSON response
        return jsonify({'prediction': prediction.tolist()})

    except ValueError as ve:
        return jsonify({"error": f"Value Error: {str(ve)}"}), 400
    except Exception as e:
        return jsonify({"error": f"An error occurred: {str(e)}"}), 500

if __name__ == '__main__':
    try:
        app.run(debug=True)
    except Exception as e:
        print(f"Error starting the Flask app: {str(e)}")

```

* Serving Flask app '__main__'

* Debug mode: on

WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.

* Running on http://127.0.0.1:5000

Press CTRL+C to quit

* Restarting with watchdog (windowsapi)

An exception has occurred, use %tb to see the full traceback.

SystemExit: 1

In [71]: *# Testing the API with curl:*
curl -X POST http://127.0.0.1:5000/predict -H "Content-Type: application/j

In [72]: *# Streamlit Deployment*

```

import streamlit as st
import joblib
import pandas as pd

# Load the trained model
try:
    model = joblib.load('final_model_xgboost.pkl')
except Exception as e:
    st.error(f"Error loading model: {str(e)}")

# Define the Streamlit app
def predict():
    st.title("Truck Delivery Prediction App")

    # Input fields for the model
    city_from = st.selectbox('City From', ['Delhi', 'Kanpur', 'Mumbai', 'Chennai'])
    city_to = st.selectbox('City To', ['Delhi', 'Kanpur', 'Mumbai', 'Chennai'])
    route_distance_km = st.number_input('Route Distance (km)', min_value=1,
    vehicle_type = st.selectbox('Vehicle Type', ['Open Truck', 'Container',
    material_weight = st.number_input('Material Weight (Tons)', min_value=0.
    truck_length_ft = st.selectbox('Truck Length (ft)', [7, 8, 14, 17, 19, 2
    truck_height_ft = st.selectbox('Truck Height (ft)', [6.0, 8.0])
    business_category = st.selectbox('Business Category', ['Manufacturer', '

```



```

weather_condition = st.selectbox('Weather Condition', ['Sunny', 'Rainy',
traffic_condition = st.selectbox('Traffic Condition', ['Light', 'Moderat
driver_rating = st.slider('Driver Rating', 1.0, 5.0, 4.5)
customer_rating = st.slider('Customer Rating', 1.0, 5.0, 4.8)

# Create a button for prediction
if st.button('Predict'):
    try:
        # Create a DataFrame with the input values
        input_data = pd.DataFrame({
            'city_from': [city_from],
            'city_to': [city_to],
            'route_distance_km': [route_distance_km],
            'vehicle_type': [vehicle_type],
            'material_weight': [material_weight],
            'truck_length_ft': [truck_length_ft],
            'truck_height_ft': [truck_height_ft],
            'business_category': [business_category],
            'weather_condition': [weather_condition],
            'traffic_condition': [traffic_condition],
            'driver_rating': [driver_rating],
            'customer_rating': [customer_rating]
        })

        # Make predictions using the loaded model
        prediction = model.predict(input_data)
        st.success(f"Predicted Price: ₹{prediction[0]:.2f}")

    except ValueError as ve:
        st.error(f"Value Error: {str(ve)}")
    except Exception as e:
        st.error(f"An error occurred: {str(e)}")

# Run the Streamlit app
if __name__ == '__main__':
    predict()

```

2024-11-01 10:22:03.385

Warning: to view this Streamlit app on a browser, run it with the following command:

```
streamlit run C:\Users\prasad jadhav\AppData\Local\Programs\Python\Python
310\lib\site-packages\ipykernel_launcher.py [ARGUMENTS]
```

In [73]: `# streamlit run your_app_name.py`

In [76]: `# Code to Optimize Multiple Models Using Optuna:`

```

import optuna
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error, accuracy_score
from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

```

```

from xgboost import XGBRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler #, LabelEncoder

# Load the dataset
# df = pd.read_csv('truck_booking_dataset.csv')

# Define the target and features
X = price_features # price_features should contain only the predictor variables
y = df['order_price']

# Label encoding for categorical features
# categorical_cols = ['city_from', 'city_to', 'vehicle_type', 'business_category']
# for col in categorical_cols:
#     le = LabelEncoder()
#     X[col] = le.fit_transform(X[col])

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the data for models that require scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Optuna objective function to optimize different models
def objective(trial):
    # Choose the model type
    model_type = trial.suggest_categorical("model", ["ridge", "decision_tree", "random_forest", "gradient_boosting", "xgboost"])

    if model_type == "ridge":
        # Ridge Regression
        alpha = trial.suggest_float("alpha", 1e-3, 10.0, log=True)
        model = Ridge(alpha=alpha)

    elif model_type == "decision_tree":
        # Decision Tree Regressor
        max_depth = trial.suggest_int("max_depth", 2, 32)
        min_samples_split = trial.suggest_int("min_samples_split", 2, 20)
        model = DecisionTreeRegressor(max_depth=max_depth, min_samples_split=min_samples_split)

    elif model_type == "random_forest":
        # Random Forest Regressor
        n_estimators = trial.suggest_int("n_estimators", 100, 1000)
        max_depth = trial.suggest_int("max_depth", 2, 32)
        model = RandomForestRegressor(n_estimators=n_estimators, max_depth=max_depth)

    elif model_type == "gradient_boosting":
        # Gradient Boosting Regressor
        learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1, log=True)
        n_estimators = trial.suggest_int("n_estimators", 100, 1000)
        max_depth = trial.suggest_int("max_depth", 2, 32)
        model = GradientBoostingRegressor(learning_rate=learning_rate, n_estimators=n_estimators, max_depth=max_depth)

    elif model_type == "xgboost":
        # XGBoost Regressor
        max_depth = trial.suggest_int("max_depth", 2, 32)
        learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1, log=True)
        n_estimators = trial.suggest_int("n_estimators", 100, 1000)
        model = XGBRegressor(max_depth=max_depth, learning_rate=learning_rate, n_estimators=n_estimators)

    # Train the model
    model.fit(X_train_scaled, y_train)

    # Evaluate the model
    y_pred = model.predict(X_test_scaled)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))

    return rmse

```

```

        learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1, log=
n_estimators = trial.suggest_int("n_estimators", 100, 1000)
max_depth = trial.suggest_int("max_depth", 2, 32)
model = XGBRegressor(learning_rate=learning_rate, n_estimators=n_est

elif model_type == "mlp":
    # Neural Network (MLP)
    hidden_layer_sizes = trial.suggest_categorical("hidden_layer_sizes",
activation = trial.suggest_categorical("activation", ["relu", "tanh"
learning_rate_init = trial.suggest_float("learning_rate_init", 1e-4,
model = MLPRegressor(hidden_layer_sizes=hidden_layer_sizes, activati

# Cross-validation scoring
score = cross_val_score(model, X_train_scaled, y_train, cv=3, scoring='n

return -score

# Create an Optuna study for hyperparameter optimization
study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=100)

# Print the best parameters and best model type
print("Best model:", study.best_trial.params)

# Train the best model on the full training set
best_model_params = study.best_trial.params
model_type = best_model_params.pop("model")

if model_type == "ridge":
    model = Ridge(**best_model_params)
elif model_type == "decision_tree":
    model = DecisionTreeRegressor(**best_model_params)
elif model_type == "random_forest":
    model = RandomForestRegressor(**best_model_params)
elif model_type == "gradient_boosting":
    model = GradientBoostingRegressor(**best_model_params)
elif model_type == "xgboost":
    model = XGBRegressor(**best_model_params)
elif model_type == "mlp":
    model = MLPRegressor(**best_model_params)

# Fit the model on the full training set
model.fit(X_train_scaled, y_train)

# Predict on the test set
y_pred = model.predict(X_test_scaled)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print(f"Test MSE: {mse}")

```

[I 2024-11-01 10:25:24,081] A new study created in memory with name: no-name-d98e71a2-aa33-466d-9229-14e44e159002

[I 2024-11-01 10:25:24,139] Trial 0 finished with value: 68663428.11762954 and parameters: {'model': 'ridge', 'alpha': 0.4161175979445562}. Best is trial 0 with value: 68663428.11762954.

[I 2024-11-01 10:25:27,445] Trial 1 finished with value: 70283473.631244 and parameters: {'model': 'xgboost', 'learning_rate': 0.004460538071699895, 'n_estimators': 171, 'max_depth': 10}. Best is trial 0 with value: 68663428.11762954.

[I 2024-11-01 10:25:27,469] Trial 2 finished with value: 68663386.20057793 and parameters: {'model': 'ridge', 'alpha': 1.9722089055681706}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:25:27,605] Trial 3 finished with value: 108750411.9074618 and parameters: {'model': 'decision_tree', 'max_depth': 30, 'min_samples_split': 11}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:26:11,235] Trial 4 finished with value: 74483055.68066157 and parameters: {'model': 'random_forest', 'n_estimators': 291, 'max_depth': 28}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:27:11,769] Trial 5 finished with value: 91219335.35091932 and parameters: {'model': 'xgboost', 'learning_rate': 0.04524494161449243, 'n_estimators': 407, 'max_depth': 23}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:27:11,799] Trial 6 finished with value: 68663425.3773238 and parameters: {'model': 'ridge', 'alpha': 0.5176822621674193}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:28:26,231] Trial 7 finished with value: 447143412.4074345 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 100), 'activation': 'tanh', 'learning_rate_init': 0.0005981755838983597}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:28:59,329] Trial 8 finished with value: 80779523.79128397 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.003200889164938601, 'n_estimators': 171, 'max_depth': 28}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:30:28,089] Trial 9 finished with value: 70686232.53790055 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (100, 100), 'activation': 'relu', 'learning_rate_init': 0.0003510009541667591}. Best is trial 2 with value: 68663386.20057793.

[I 2024-11-01 10:30:28,135] Trial 10 finished with value: 68663239.51365662 and parameters: {'model': 'ridge', 'alpha': 7.4604331373263815}. Best is trial 10 with value: 68663239.51365662.

[I 2024-11-01 10:30:28,181] Trial 11 finished with value: 68663197.8349422 and parameters: {'model': 'ridge', 'alpha': 9.032128043222931}. Best is trial 11 with value: 68663197.8349422.

[I 2024-11-01 10:30:28,214] Trial 12 finished with value: 68663182.00951947 and parameters: {'model': 'ridge', 'alpha': 9.630345653458786}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:30:28,267] Trial 13 finished with value: 68663439.15099548 and parameters: {'model': 'ridge', 'alpha': 0.00741650100015105}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:30:47,297] Trial 14 finished with value: 71649207.14858495 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.09221526022065295, 'n_estimators': 988, 'max_depth': 2}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:30:47,435] Trial 15 finished with value: 90155355.30729155 and parameters: {'model': 'decision_tree', 'max_depth': 15, 'min_samples_split': 20}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:30:47,481] Trial 16 finished with value: 68663202.36706114 and parameters: {'model': 'ridge', 'alpha': 8.860955974262941}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:31:04,929] Trial 17 finished with value: 68754764.53153673 and parameters: {'model': 'random_forest', 'n_estimators': 761, 'max_depth': 3}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:31:04,966] Trial 18 finished with value: 68663438.73461847 and parameters: {'model': 'ridge', 'alpha': 0.02283333148843314}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:31:05,004] Trial 19 finished with value: 68663439.32132785 and parameters: {'model': 'ridge', 'alpha': 0.0011099043745178626}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:31:57,254] Trial 20 finished with value: 375510295.734067 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 50), 'activation': 'tanh', 'learning_rate_init': 0.008383116576355697}. Best is trial 12 with value: 68663182.00951947.

[I 2024-11-01 10:31:57,308] Trial 21 finished with value: 68663172.4173289 and parameters: {'model': 'ridge', 'alpha': 9.993329168835508}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:31:57,354] Trial 22 finished with value: 68663204.93295261 and parameters: {'model': 'ridge', 'alpha': 8.764074621023136}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:31:57,446] Trial 23 finished with value: 68663409.46826835 and parameters: {'model': 'ridge', 'alpha': 1.1077763342027287}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:31:57,527] Trial 24 finished with value: 68663362.68649116 and parameters: {'model': 'ridge', 'alpha': 2.847484363675904}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:33:30,378] Trial 25 finished with value: 74235026.76154555 and parameters: {'model': 'xgboost', 'learning_rate': 0.0010842199043400134, 'n_estimators': 633, 'max_depth': 19}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:34:21,692] Trial 26 finished with value: 70279978.13412084 and parameters: {'model': 'random_forest', 'n_estimators': 818, 'max_depth': 10}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:51,740] Trial 27 finished with value: 117580609.75659454 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.02155304124740829, 'n_estimators': 480, 'max_depth': 22}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:51,857] Trial 28 finished with value: 81616099.29211701 and parameters: {'model': 'decision_tree', 'max_depth': 9, 'min_samples_split': 3}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:51,905] Trial 29 finished with value: 68663435.94352134 and parameters: {'model': 'ridge', 'alpha': 0.12619047702739875}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:51,943] Trial 30 finished with value: 68663354.68117605 and parameters: {'model': 'ridge', 'alpha': 3.145858507648962}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:51,981] Trial 31 finished with value: 68663204.80646567 and parameters: {'model': 'ridge', 'alpha': 8.76884994698507}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:52,017] Trial 32 finished with value: 68663263.18795009 and parameters: {'model': 'ridge', 'alpha': 6.570124331041535}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:35:52,058] Trial 33 finished with value: 68663360.39689736 and parameters: {'model': 'ridge', 'alpha': 2.9328019237659544}. Best is trial

21 with value: 68663172.4173289.
[I 2024-11-01 10:35:52,096] Trial 34 finished with value: 68663193.9039505 and parameters: {'model': 'ridge', 'alpha': 9.18064923146643}. Best is trial 21 with value: 68663172.4173289.
[I 2024-11-01 10:36:29,589] Trial 35 finished with value: 86181072.42693734 and parameters: {'model': 'xgboost', 'learning_rate': 0.011264266172384322, 'n_estimators': 991, 'max_depth': 13}. Best is trial 21 with value: 68663172.4173289.
[I 2024-11-01 10:36:29,632] Trial 36 finished with value: 68663412.04716446 and parameters: {'model': 'ridge', 'alpha': 1.0120681728682561}. Best is trial 21 with value: 68663172.4173289.
[I 2024-11-01 10:36:29,765] Trial 37 finished with value: 94194875.37447363 and parameters: {'model': 'decision_tree', 'max_depth': 32, 'min_samples_split': 20}. Best is trial 21 with value: 68663172.4173289.
[I 2024-11-01 10:36:52,076] Trial 38 finished with value: 68965976.72235405 and parameters: {'model': 'random_forest', 'n_estimators': 624, 'max_depth': 5}. Best is trial 21 with value: 68663172.4173289.
[I 2024-11-01 10:36:52,130] Trial 39 finished with value: 68663341.20960145 and parameters: {'model': 'ridge', 'alpha': 3.648418023356419}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:37:53,478] Trial 40 finished with value: 331672719.26505977 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 100), 'activation': 'relu', 'learning_rate_init': 0.00010459811575120866}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:37:53,523] Trial 41 finished with value: 68663200.39053816 and parameters: {'model': 'ridge', 'alpha': 8.935598576306964}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:37:53,563] Trial 42 finished with value: 68663176.23936951 and parameters: {'model': 'ridge', 'alpha': 9.848661964720929}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:37:53,603] Trial 43 finished with value: 68663324.30895157 and parameters: {'model': 'ridge', 'alpha': 4.279694590967343}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:37:53,639] Trial 44 finished with value: 68663399.53962703 and parameters: {'model': 'ridge', 'alpha': 1.47643864612499}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:39:08,018] Trial 45 finished with value: 74858885.09051067 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.001047984785929564, 'n_estimators': 349, 'max_depth': 24}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:41:07,096] Trial 46 finished with value: 88022859.70089124 and parameters: {'model': 'xgboost', 'learning_rate': 0.004177727654344027, 'n_estimators': 811, 'max_depth': 18}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:41:07,137] Trial 47 finished with value: 68663327.4942287 and parameters: {'model': 'ridge', 'alpha': 4.16064970521638}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:41:07,179] Trial 48 finished with value: 68663180.15557887 and parameters: {'model': 'ridge', 'alpha': 9.700478777027246}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:41:07,220] Trial 49 finished with value: 68663318.1245788 and parameters: {'model': 'ridge', 'alpha': 4.510916109929547}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:50,900] Trial 50 finished with value: 321605198.13700265 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (100, 100), 'activation': 'tanh', 'learning_rate_init': 0.007320204607182357}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:50,936] Trial 51 finished with value: 68663172.81617485 and parameters: {'model': 'ridge', 'alpha': 9.978230359395363}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:50,981] Trial 52 finished with value: 68663190.04134978 and parameters: {'model': 'ridge', 'alpha': 9.32663435882943}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:51,018] Trial 53 finished with value: 68663311.05879058 and parameters: {'model': 'ridge', 'alpha': 4.775237573421291}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:51,066] Trial 54 finished with value: 68663380.99942403 and parameters: {'model': 'ridge', 'alpha': 2.1656673157052944}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:51,146] Trial 55 finished with value: 70467598.43991484 and parameters: {'model': 'decision_tree', 'max_depth': 5, 'min_samples_split': 2}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:42:51,179] Trial 56 finished with value: 68663306.93011063 and parameters: {'model': 'ridge', 'alpha': 4.929757771205937}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:43:04,359] Trial 57 finished with value: 74639003.82406723 and parameters: {'model': 'ridge', 'alpha': 1.47643864612499}. Best is trial 21 with value: 68663172.4173289.

nd parameters: {'model': 'random_forest', 'n_estimators': 114, 'max_depth': 26}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:43:04,406] Trial 58 finished with value: 68663384.30286078 and parameters: {'model': 'ridge', 'alpha': 2.0427854096082965}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,708] Trial 59 finished with value: 112164899.2254107 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.0208337381319872, 'n_estimators': 548, 'max_depth': 21}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,755] Trial 60 finished with value: 68663422.98343092 and parameters: {'model': 'ridge', 'alpha': 0.6064264712563425}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,801] Trial 61 finished with value: 68663201.57490735 and parameters: {'model': 'ridge', 'alpha': 8.890869855545956}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,833] Trial 62 finished with value: 68663292.4542823 and parameters: {'model': 'ridge', 'alpha': 5.471950722456521}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,882] Trial 63 finished with value: 68663180.17649023 and parameters: {'model': 'ridge', 'alpha': 9.699687655532726}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,930] Trial 64 finished with value: 68663433.76283385 and parameters: {'model': 'ridge', 'alpha': 0.20696003067822813}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:42,985] Trial 65 finished with value: 68663275.93168502 and parameters: {'model': 'ridge', 'alpha': 6.091605657378613}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:52,721] Trial 66 finished with value: 88156609.23803066 and parameters: {'model': 'xgboost', 'learning_rate': 0.08621203286088601, 'n_estimators': 250, 'max_depth': 13}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:44:52,768] Trial 67 finished with value: 68663438.19897753 and parameters: {'model': 'ridge', 'alpha': 0.04266681891898758}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,092] Trial 68 finished with value: 69786952.87031387 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 50), 'activation': 'relu', 'learning_rate_init': 0.002634158937293017}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,140] Trial 69 finished with value: 68663280.76578356 and parameters: {'model': 'ridge', 'alpha': 5.910221712939825}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,184] Trial 70 finished with value: 68663258.72906186 and parameters: {'model': 'ridge', 'alpha': 6.737672991333914}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,233] Trial 71 finished with value: 68663239.13354975 and parameters: {'model': 'ridge', 'alpha': 7.474742052536655}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,276] Trial 72 finished with value: 68663192.20251927 and parameters: {'model': 'ridge', 'alpha': 9.244948150254551}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,328] Trial 73 finished with value: 68663189.22902602 and parameters: {'model': 'ridge', 'alpha': 9.357341788760497}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,382] Trial 74 finished with value: 68663183.66536163 and parameters: {'model': 'ridge', 'alpha': 9.567715707450418}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:45:47,495] Trial 75 finished with value: 73922826.24768418 and parameters: {'model': 'decision_tree', 'max_depth': 7, 'min_samples_split': 11}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:47:34,024] Trial 76 finished with value: 73044488.19842939 and parameters: {'model': 'random_forest', 'n_estimators': 875, 'max_depth': 16}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:47:34,063] Trial 77 finished with value: 68663362.6934464 and parameters: {'model': 'ridge', 'alpha': 2.8472252146974903}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:19,847] Trial 78 finished with value: 76912734.52317987 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.002156541791947052, 'n_estimators': 707, 'max_depth': 13}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:19,882] Trial 79 finished with value: 68663283.67910737 and parameters: {'model': 'ridge', 'alpha': 5.800943979334838}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:19,916] Trial 80 finished with value: 68663439.24001162 and parameters: {'model': 'ridge', 'alpha': 0.004120645592304338}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:19,959] Trial 81 finished with value: 68663177.9432019 and parameters: {'model': 'ridge', 'alpha': 9.78418561826308}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:19,998] Trial 82 finished with value: 68663176.17035209 and parameters: {'model': 'ridge', 'alpha': 9.85127391421915}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:20,041] Trial 83 finished with value: 68663175.0925818 and parameters: {'model': 'ridge', 'alpha': 9.892063870761616}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:20,082] Trial 84 finished with value: 68663280.19485657 and parameters: {'model': 'ridge', 'alpha': 5.931640095245019}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:49:20,119] Trial 85 finished with value: 68663347.02712584 and parameters: {'model': 'ridge', 'alpha': 3.431325380658104}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:50:43,899] Trial 86 finished with value: 90032087.15485202 and parameters: {'model': 'xgboost', 'learning_rate': 0.008632417827405066, 'n_estimators': 477, 'max_depth': 26}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:50:43,943] Trial 87 finished with value: 68663277.00856246 and parameters: {'model': 'ridge', 'alpha': 6.051192970005129}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:50:43,985] Trial 88 finished with value: 68663337.41720887 and parameters: {'model': 'ridge', 'alpha': 3.7899952030713893}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,682] Trial 89 finished with value: 457368078.3260831 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 100), 'activation': 'tanh', 'learning_rate_init': 0.0001278985892255755}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,720] Trial 90 finished with value: 68663273.24849682 and parameters: {'model': 'ridge', 'alpha': 6.192315233273207}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,758] Trial 91 finished with value: 68663186.95080769 and parameters: {'model': 'ridge', 'alpha': 9.443474117044277}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,789] Trial 92 finished with value: 68663183.91001399 and parameters: {'model': 'ridge', 'alpha': 9.558462810645214}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,822] Trial 93 finished with value: 68663172.59836781 and parameters: {'model': 'ridge', 'alpha': 9.986475652231903}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,871] Trial 94 finished with value: 68663253.33098139 and parameters: {'model': 'ridge', 'alpha': 6.94059669266395}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:52:04,902] Trial 95 finished with value: 68663321.15183832 and parameters: {'model': 'ridge', 'alpha': 4.397717988601275}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:54:34,314] Trial 96 finished with value: 74366920.8828585 and parameters: {'model': 'random_forest', 'n_estimators': 909, 'max_depth': 32}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:54:34,506] Trial 97 finished with value: 97986079.92063189 and parameters: {'model': 'decision_tree', 'max_depth': 20, 'min_samples_split': 16}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:54:34,557] Trial 98 finished with value: 68663364.21888255 and parameters: {'model': 'ridge', 'alpha': 2.7903916094161874}. Best is trial 21 with value: 68663172.4173289.

[I 2024-11-01 10:54:34,592] Trial 99 finished with value: 68663242.07998374 and parameters: {'model': 'ridge', 'alpha': 7.363837080302506}. Best is trial 21 with value: 68663172.4173289.

Best model: {'model': 'ridge', 'alpha': 9.993329168835508}

Test MSE: 69594019.82830097

```
In [77]: import optuna
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from xgboost import XGBRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler
import pickle

# Define the target and features
X = price_features # price_features should contain only the predictor variables
y = df['order_price']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the data for models that require scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Optuna objective function to optimize different models
def objective(trial):
    # Choose the model type
    model_type = trial.suggest_categorical("model", ["ridge", "decision_tree", "random_forest", "gradient_boosting"])

    if model_type == "ridge":
        # Ridge Regression
        alpha = trial.suggest_float("alpha", 1e-3, 10.0, log=True)
        model = Ridge(alpha=alpha)

    elif model_type == "decision_tree":
        # Decision Tree Regressor
        max_depth = trial.suggest_int("max_depth", 2, 32)
        min_samples_split = trial.suggest_int("min_samples_split", 2, 20)
        model = DecisionTreeRegressor(max_depth=max_depth, min_samples_split=min_samples_split)

    elif model_type == "random_forest":
        # Random Forest Regressor
        n_estimators = trial.suggest_int("n_estimators", 100, 1000)
        max_depth = trial.suggest_int("max_depth", 2, 32)
        model = RandomForestRegressor(n_estimators=n_estimators, max_depth=max_depth)

    elif model_type == "gradient_boosting":
        # Gradient Boosting Regressor
        learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1, log=True)
        n_estimators = trial.suggest_int("n_estimators", 100, 1000)
        max_depth = trial.suggest_int("max_depth", 2, 32)
        model = GradientBoostingRegressor(learning_rate=learning_rate, n_estimators=n_estimators, max_depth=max_depth)
```

```

        model = GradientBoostingRegressor(learning_rate=learning_rate, n_est

elif model_type == "xgboost":
    # XGBoost Regressor
    learning_rate = trial.suggest_float("learning_rate", 1e-3, 0.1, log=
    n_estimators = trial.suggest_int("n_estimators", 100, 1000)
    max_depth = trial.suggest_int("max_depth", 2, 32)
    model = XGBRegressor(learning_rate=learning_rate, n_estimators=n_est

elif model_type == "mlp":
    # Neural Network (MLP)
    hidden_layer_sizes = trial.suggest_categorical("hidden_layer_sizes",
    activation = trial.suggest_categorical("activation", ["relu", "tanh"
    learning_rate_init = trial.suggest_float("learning_rate_init", 1e-4,
    model = MLPRegressor(hidden_layer_sizes=hidden_layer_sizes, activati

# Cross-validation scoring
score = cross_val_score(model, X_train_scaled, y_train, cv=3, scoring='n

return -score

# Create an Optuna study for hyperparameter optimization
study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=100)

# Print the best parameters and best model type
print("Best model parameters:", study.best_trial.params)

# Train the best model on the full training set
best_model_params = study.best_trial.params
model_type = best_model_params.pop("model")

# Define the best model using the best parameters
if model_type == "ridge":
    best_model = Ridge(**best_model_params)
elif model_type == "decision_tree":
    best_model = DecisionTreeRegressor(**best_model_params)
elif model_type == "random_forest":
    best_model = RandomForestRegressor(**best_model_params)
elif model_type == "gradient_boosting":
    best_model = GradientBoostingRegressor(**best_model_params)
elif model_type == "xgboost":
    best_model = XGBRegressor(**best_model_params)
elif model_type == "mlp":
    best_model = MLPRegressor(**best_model_params)

# Fit the best model on the full training set
best_model.fit(X_train_scaled, y_train)

# Save the best model as a .pkl file
pickle.dump(best_model, 'best_model.pkl', 'wb')
# print("Best model saved as 'best_model.pkl'")

```

[I 2024-11-01 10:59:33,971] A new study created in memory with name: no-name-ee069f38-6f6b-4a74-80ec-60ea99f0ad5f

[I 2024-11-01 11:00:31,375] Trial 0 finished with value: 112131791.75526951 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.03546117757472154, 'n_estimators': 278, 'max_depth': 21}. Best is trial 0 with value: 112131791.75526951.

[I 2024-11-01 11:01:40,648] Trial 1 finished with value: 73926164.257764 and parameters: {'model': 'random_forest', 'n_estimators': 444, 'max_depth': 20}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:03:10,566] Trial 2 finished with value: 91020685.58255081 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (100, 100), 'activation': 'relu', 'learning_rate_init': 0.0001625450090123917}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:03:43,593] Trial 3 finished with value: 90674104.99670553 and parameters: {'model': 'xgboost', 'learning_rate': 0.03553147724738988, 'n_estimators': 196, 'max_depth': 23}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:06:41,352] Trial 4 finished with value: 91653761.8219328 and parameters: {'model': 'xgboost', 'learning_rate': 0.007402053426589449, 'n_estimators': 934, 'max_depth': 31}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:08:02,658] Trial 5 finished with value: 439621011.68819284 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 100), 'activation': 'tanh', 'learning_rate_init': 0.0009533560316975296}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:08:18,910] Trial 6 finished with value: 82569769.00903039 and parameters: {'model': 'xgboost', 'learning_rate': 0.013702583608935476, 'n_estimators': 241, 'max_depth': 14}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:08:19,007] Trial 7 finished with value: 78595527.29848376 and parameters: {'model': 'decision_tree', 'max_depth': 9, 'min_samples_split': 19}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:09:07,008] Trial 8 finished with value: 85423377.56786168 and parameters: {'model': 'xgboost', 'learning_rate': 0.007265910183203576, 'n_estimators': 866, 'max_depth': 14}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:09:07,170] Trial 9 finished with value: 98950063.68894637 and parameters: {'model': 'decision_tree', 'max_depth': 13, 'min_samples_split': 8}. Best is trial 1 with value: 73926164.257764.

[I 2024-11-01 11:09:30,441] Trial 10 finished with value: 68937098.82760148 and parameters: {'model': 'random_forest', 'n_estimators': 555, 'max_depth': 5}. Best is trial 10 with value: 68937098.82760148.

[I 2024-11-01 11:09:49,397] Trial 11 finished with value: 68829507.82523823 and parameters: {'model': 'random_forest', 'n_estimators': 530, 'max_depth': 4}. Best is trial 11 with value: 68829507.82523823.

[I 2024-11-01 11:10:04,335] Trial 12 finished with value: 68685605.69353737 and parameters: {'model': 'random_forest', 'n_estimators': 639, 'max_depth': 2}. Best is trial 12 with value: 68685605.69353737.

[I 2024-11-01 11:10:04,384] Trial 13 finished with value: 68663326.62438862 and parameters: {'model': 'ridge', 'alpha': 4.193155528525295}. Best is trial 13 with value: 68663326.62438862.

[I 2024-11-01 11:10:04,444] Trial 14 finished with value: 68663281.28098387 and parameters: {'model': 'ridge', 'alpha': 5.890894800544558}. Best is trial 14 with value: 68663281.28098387.

[I 2024-11-01 11:10:04,500] Trial 15 finished with value: 68663228.91314222 and parameters: {'model': 'ridge', 'alpha': 7.859654740136146}. Best is trial

15 with value: 68663228.91314222.
[I 2024-11-01 11:10:04,548] Trial 16 finished with value: 68663258.35382725 and parameters: {'model': 'ridge', 'alpha': 6.751775781663592}. Best is trial 15 with value: 68663228.91314222.
[I 2024-11-01 11:10:04,605] Trial 17 finished with value: 68663198.66260736 and parameters: {'model': 'ridge', 'alpha': 9.000863364753389}. Best is trial 17 with value: 68663198.66260736.
[I 2024-11-01 11:10:04,662] Trial 18 finished with value: 68663439.28329735 and parameters: {'model': 'ridge', 'alpha': 0.002517984937074765}. Best is trial 17 with value: 68663198.66260736.
[I 2024-11-01 11:10:04,708] Trial 19 finished with value: 68663433.7354982 and parameters: {'model': 'ridge', 'alpha': 0.20797259589413097}. Best is trial 17 with value: 68663198.66260736.
[I 2024-11-01 11:13:18,190] Trial 20 finished with value: 89547803.9417217 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.0010416368443338197, 'n_estimators': 780, 'max_depth': 31}. Best is trial 17 with value: 68663198.66260736.
[I 2024-11-01 11:13:18,233] Trial 21 finished with value: 68663187.63629206 and parameters: {'model': 'ridge', 'alpha': 9.417556352371095}. Best is trial 21 with value: 68663187.63629206.
[I 2024-11-01 11:13:18,279] Trial 22 finished with value: 68663175.29959507 and parameters: {'model': 'ridge', 'alpha': 9.884228831525226}. Best is trial 22 with value: 68663175.29959507.
[I 2024-11-01 11:13:18,318] Trial 23 finished with value: 68663425.15387116 and parameters: {'model': 'ridge', 'alpha': 0.5259651553580619}. Best is trial 22 with value: 68663175.29959507.
[I 2024-11-01 11:13:18,358] Trial 24 finished with value: 68663411.70680614 and parameters: {'model': 'ridge', 'alpha': 1.0246984088127367}. Best is trial 22 with value: 68663175.29959507.
[I 2024-11-01 11:13:18,399] Trial 25 finished with value: 68663439.15953805 and parameters: {'model': 'ridge', 'alpha': 0.0071002080373900696}. Best is trial 22 with value: 68663175.29959507.
[I 2024-11-01 11:13:18,436] Trial 26 finished with value: 68663400.7782608 and parameters: {'model': 'ridge', 'alpha': 1.430430221479741}. Best is trial 22 with value: 68663175.29959507.
[I 2024-11-01 11:13:18,475] Trial 27 finished with value: 68663174.89673787 and parameters: {'model': 'ridge', 'alpha': 9.899476298425093}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:13:59,893] Trial 28 finished with value: 69698177.55892675 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 50), 'activation': 'relu', 'learning_rate_init': 0.008998111703920574}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:14:26,187] Trial 29 finished with value: 70213313.39012311 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.0012546538517486736, 'n_estimators': 117, 'max_depth': 27}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:14:26,410] Trial 30 finished with value: 142390031.3780454 and parameters: {'model': 'decision_tree', 'max_depth': 26, 'min_samples_split': 2}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:14:26,447] Trial 31 finished with value: 68663186.31992017 and parameters: {'model': 'ridge', 'alpha': 9.467328930608}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:14:26,476] Trial 32 finished with value: 68663390.48541798 and parameters: {'model': 'ridge', 'alpha': 1.812895458047031}. Best is trial 27 with value: 68663174.89673787.
[I 2024-11-01 11:14:26,523] Trial 33 finished with value: 68663178.50488697 a

nd parameters: {'model': 'ridge', 'alpha': 9.762932387863508}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:14:55,700] Trial 34 finished with value: 70100413.61892541 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.002979747069716516, 'n_estimators': 392, 'max_depth': 8}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:14:55,746] Trial 35 finished with value: 68663381.95354599 and parameters: {'model': 'ridge', 'alpha': 2.130172251069295}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:16:15,779] Trial 36 finished with value: 299667035.41928643 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 100), 'activation': 'tanh', 'learning_rate_init': 0.008696218991372793}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:16:15,819] Trial 37 finished with value: 68663438.62875076 and parameters: {'model': 'ridge', 'alpha': 0.026753287247137946}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:16:15,857] Trial 38 finished with value: 68663368.7463354 and parameters: {'model': 'ridge', 'alpha': 2.621753256139958}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:17:17,713] Trial 39 finished with value: 89942805.6074364 and parameters: {'model': 'xgboost', 'learning_rate': 0.08495524965717507, 'n_estimators': 686, 'max_depth': 17}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:17:17,840] Trial 40 finished with value: 81611605.49670173 and parameters: {'model': 'decision_tree', 'max_depth': 10, 'min_samples_split': 18}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:17:17,880] Trial 41 finished with value: 68663185.62875855 and parameters: {'model': 'ridge', 'alpha': 9.49346425378212}. Best is trial 27 with value: 68663174.89673787.

[I 2024-11-01 11:17:17,927] Trial 42 finished with value: 68663174.37925433 and parameters: {'model': 'ridge', 'alpha': 9.919062941112978}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:17:17,973] Trial 43 finished with value: 68663345.03537782 and parameters: {'model': 'ridge', 'alpha': 3.5056396822374385}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:18:49,891] Trial 44 finished with value: 108854886.76997614 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (100, 100), 'activation': 'relu', 'learning_rate_init': 0.00014477675902466267}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:21:30,463] Trial 45 finished with value: 74285111.55797626 and parameters: {'model': 'random_forest', 'n_estimators': 966, 'max_depth': 27}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:22:16,401] Trial 46 finished with value: 77355441.188477 and parameters: {'model': 'xgboost', 'learning_rate': 0.0028366704544452483, 'n_estimators': 370, 'max_depth': 18}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:22:16,464] Trial 47 finished with value: 68663334.13179599 and parameters: {'model': 'ridge', 'alpha': 3.9126818927669604}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:22:16,512] Trial 48 finished with value: 68663177.46812922 and parameters: {'model': 'ridge', 'alpha': 9.80216236523409}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:22:16,565] Trial 49 finished with value: 68663419.68861735 and parameters: {'model': 'ridge', 'alpha': 0.7285973724861334}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:27,724] Trial 50 finished with value: 74278734.28929675 and parameters: {'model': 'random_forest', 'n_estimators': 784, 'max_depth': 32}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:27,777] Trial 51 finished with value: 68663192.42292733 and parameters: {'model': 'ridge', 'alpha': 9.2366181718989}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:27,833] Trial 52 finished with value: 68663335.07308923 and parameters: {'model': 'ridge', 'alpha': 3.8775278988253854}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:27,882] Trial 53 finished with value: 68663315.64682506 and parameters: {'model': 'ridge', 'alpha': 4.6035878018735215}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:27,944] Trial 54 finished with value: 68663190.83166659 and parameters: {'model': 'ridge', 'alpha': 9.296760852351381}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:28,134] Trial 55 finished with value: 108485618.36304604 and parameters: {'model': 'decision_tree', 'max_depth': 24, 'min_samples_split': 11}. Best is trial 42 with value: 68663174.37925433.

[I 2024-11-01 11:24:28,182] Trial 56 finished with value: 68663358.29708788 and parameters: {'model': 'ridge', 'alpha': 3.011061733234698}. Best is trial

42 with value: 68663174.37925433.
[I 2024-11-01 11:24:41,454] Trial 57 finished with value: 83547932.46947806 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.09518075254772504, 'n_estimators': 121, 'max_depth': 11}. Best is trial 42 with value: 68663174.37925433.
[I 2024-11-01 11:24:41,523] Trial 58 finished with value: 68663437.4447334 and parameters: {'model': 'ridge', 'alpha': 0.0705961327605018}. Best is trial 42 with value: 68663174.37925433.
[I 2024-11-01 11:24:44,210] Trial 59 finished with value: 73154585.47437084 and parameters: {'model': 'xgboost', 'learning_rate': 0.01747375772181485, 'n_estimators': 527, 'max_depth': 6}. Best is trial 42 with value: 68663174.37925433.
[I 2024-11-01 11:24:44,275] Trial 60 finished with value: 68663293.31926104 and parameters: {'model': 'ridge', 'alpha': 5.439534523153324}. Best is trial 42 with value: 68663174.37925433.
[I 2024-11-01 11:24:44,338] Trial 61 finished with value: 68663172.73313762 and parameters: {'model': 'ridge', 'alpha': 9.981373794788786}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:24:44,404] Trial 62 finished with value: 68663315.50653814 and parameters: {'model': 'ridge', 'alpha': 4.6088353135013085}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:24:44,452] Trial 63 finished with value: 68663290.8689352 and parameters: {'model': 'ridge', 'alpha': 5.531369727467178}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:24:44,509] Trial 64 finished with value: 68663177.78757979 and parameters: {'model': 'ridge', 'alpha': 9.790074278004049}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:24:44,559] Trial 65 finished with value: 68663293.33412008 and parameters: {'model': 'ridge', 'alpha': 5.438977682007707}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:24:44,606] Trial 66 finished with value: 68663282.09293216 and parameters: {'model': 'ridge', 'alpha': 5.860437548966841}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:25:51,125] Trial 67 finished with value: 445410800.2677758 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (50, 50), 'activation': 'tanh', 'learning_rate_init': 0.0013768945306068395}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:25:51,170] Trial 68 finished with value: 68663174.38910888 and parameters: {'model': 'ridge', 'alpha': 9.91868993939105}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,304] Trial 69 finished with value: 74351048.20409407 and parameters: {'model': 'random_forest', 'n_estimators': 706, 'max_depth': 29}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,352] Trial 70 finished with value: 68663288.15204097 and parameters: {'model': 'ridge', 'alpha': 5.633217523550235}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,416] Trial 71 finished with value: 68663175.1574439 and parameters: {'model': 'ridge', 'alpha': 9.889608955454245}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,465] Trial 72 finished with value: 68663262.2950107 and parameters: {'model': 'ridge', 'alpha': 6.603672714168007}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,526] Trial 73 finished with value: 68663366.37530492 and parameters: {'model': 'ridge', 'alpha': 2.7100614111436814}. Best is trial 61 with value: 68663172.73313762.
[I 2024-11-01 11:27:46,558] Trial 74 finished with value: 68663269.3475378 and

d parameters: {'model': 'ridge', 'alpha': 6.3387723359618295}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:26,766] Trial 75 finished with value: 89019247.82036363 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.003699630388100432, 'n_estimators': 852, 'max_depth': 16}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:26,864] Trial 76 finished with value: 68663439.3239317 and parameters: {'model': 'ridge', 'alpha': 0.0010134965549025556}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:26,960] Trial 77 finished with value: 68851525.15931912 and parameters: {'model': 'decision_tree', 'max_depth': 2, 'min_samples_split': 2}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,007] Trial 78 finished with value: 68663185.8259406 and parameters: {'model': 'ridge', 'alpha': 9.48600793171511}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,070] Trial 79 finished with value: 68663430.19996034 and parameters: {'model': 'ridge', 'alpha': 0.3389549308785647}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,121] Trial 80 finished with value: 68663259.83238366 and parameters: {'model': 'ridge', 'alpha': 6.696208388613386}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,179] Trial 81 finished with value: 68663179.84050132 and parameters: {'model': 'ridge', 'alpha': 9.712399004011218}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,229] Trial 82 finished with value: 68663175.29283471 and parameters: {'model': 'ridge', 'alpha': 9.884484695120214}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,290] Trial 83 finished with value: 68663337.34738299 and parameters: {'model': 'ridge', 'alpha': 3.7926023522392978}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,339] Trial 84 finished with value: 68663239.58409439 and parameters: {'model': 'ridge', 'alpha': 7.457781596249952}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:30:27,387] Trial 85 finished with value: 68663259.0809636 and parameters: {'model': 'ridge', 'alpha': 6.724447544442034}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:31:53,148] Trial 86 finished with value: 91011559.58907944 and parameters: {'model': 'xgboost', 'learning_rate': 0.036033153750373526, 'n_estimators': 610, 'max_depth': 20}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,451] Trial 87 finished with value: 70053906.65758651 and parameters: {'model': 'mlp', 'hidden_layer_sizes': (100, 100), 'activation': 'relu', 'learning_rate_init': 0.0008913320404642503}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,513] Trial 88 finished with value: 68663398.64904138 and parameters: {'model': 'ridge', 'alpha': 1.5095219037751253}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,571] Trial 89 finished with value: 68663175.3619619 and parameters: {'model': 'ridge', 'alpha': 9.881868398303055}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,607] Trial 90 finished with value: 68663373.3667783 and parameters: {'model': 'ridge', 'alpha': 2.4497163137202813}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,656] Trial 91 finished with value: 68663242.30499963 and parameters: {'model': 'ridge', 'alpha': 7.355368516355657}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,719] Trial 92 finished with value: 68663173.84165607 and parameters: {'model': 'ridge', 'alpha': 9.93941182592321}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,766] Trial 93 finished with value: 68663323.42011787 and parameters: {'model': 'ridge', 'alpha': 4.312919020257117}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:33:17,830] Trial 94 finished with value: 68663246.95398425 and parameters: {'model': 'ridge', 'alpha': 7.180437815460158}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:34:09,082] Trial 95 finished with value: 74365145.90719002 and parameters: {'model': 'random_forest', 'n_estimators': 318, 'max_depth': 24}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:34:09,145] Trial 96 finished with value: 68663438.94410203 and parameters: {'model': 'ridge', 'alpha': 0.015076899643281646}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:35:57,883] Trial 97 finished with value: 88788938.93956111 and parameters: {'model': 'gradient_boosting', 'learning_rate': 0.0019439160043005915, 'n_estimators': 470, 'max_depth': 22}. Best is trial 61 with value:

68663172.73313762.

[I 2024-11-01 11:35:58,055] Trial 98 finished with value: 98593193.53870712 and parameters: {'model': 'decision_tree', 'max_depth': 19, 'min_samples_split': 15}. Best is trial 61 with value: 68663172.73313762.

[I 2024-11-01 11:35:58,113] Trial 99 finished with value: 68663246.01176262 and parameters: {'model': 'ridge', 'alpha': 7.215885965570791}. Best is trial 61 with value: 68663172.73313762.

Best model parameters: {'model': 'ridge', 'alpha': 9.981373794788786}

TypeError

Traceback (most recent call last)

Cell In[77], line 101

```
98 best_model.fit(X_train_scaled, y_train)
100 # Save the best model as a .pkl file
--> 101 pickle.dump(best_model, 'best_model.pkl', 'wb')
102 print("Best model saved as 'best_model.pkl'")
```

TypeError: 'str' object cannot be interpreted as an integer

```
In [79]: pickle.dump(best_model, open('base_model.pkl', 'wb'))
        model = pickle.load(open('base_model.pkl', 'rb'))
```

```
In [81]: price_features.head()
```

```
Out[81]:
```

	city_from	city_to	vehicle_type	material_weight	truck_length_ft	truck_height_ft	business_category
0	0	3	0	1.61	8	6.0	
1	0	4	0	1.98	20	8.0	
2	1	2	0	7.28	7	6.0	
3	2	0	0	8.69	19	8.0	
4	2	1	1	7.60	14	6.0	

```
In [83]: price_features.dtypes
```

```
Out[83]: city_from      int64
city_to      int64
vehicle_type  int64
material_weight  float64
truck_length_ft  int64
truck_height_ft  float64
business_category  int64
dtype: object
```

```
In [84]: price_features.head()
```

```
Out[84]:
```

	city_from	city_to	vehicle_type	material_weight	truck_length_ft	truck_height_ft	business_category
0	0	3	0	1.61	8	6.0	
1	0	4	0	1.98	20	8.0	
2	1	2	0	7.28	7	6.0	
3	2	0	0	8.69	19	8.0	
4	2	1	1	7.60	14	6.0	

```
In [86]: y.head()
```

```
Out[86]: 0    21282.00
         1    32028.75
         2    20158.93
         3    21944.58
         4    14546.95
         Name: order_price, dtype: float64
```

```
In [88]: print('Enter Truck Details')

city_from = int(input('Enter the City From: '))
city_to = int(input('Enter the City To: '))
vehicle_type = int(input('Enter the Vehicle Type: '))
material_weight = float(input('Enter the Material Weight: '))
truck_length_ft = int(input('Enter the Truck Length ft: '))
truck_height_ft = float(input('Enter the Truck Height ft: '))
business_category = int(input('Enter the Business Category: '))

input_point = np.array([[city_from, city_to, vehicle_type, material_weight,
                           truck_length_ft, truck_height_ft, business_category]])

model.predict(input_point)
```

```
Enter Truck Details
Enter the City From: 0
Enter the City To: 3
Enter the Vehicle Type: 0
Enter the Material Weight: 1.61
Enter the Truck Length ft: 8
Enter the Truck Height ft: 6.0
Enter the Business Category: 0
```

```
Out[88]: array([21259.9722974])
```

```
In [89]: df.head()
```

```
Out[89]:
```

	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight	tru
0	ORDI000	CVS9726	0	3	1144	0	1.61	
1	ORDI001	CVS7464	0	4	1912	0	1.98	
2	ORDI002	CVS4668	1	2	1145	0	7.28	
3	ORDI003	CVS2876	2	0	1247	0	8.69	
4	ORDI004	CVS4137	2	1	569	1	7.60	

```
In [90]: price_features.head()
```

```
Out[90]:
```

	city_from	city_to	vehicle_type	material_weight	truck_length_ft	truck_height_ft	business_cate
0	0	3	0	1.61	8	6.0	
1	0	4	0	1.98	20	8.0	
2	1	2	0	7.28	7	6.0	
3	2	0	0	8.69	19	8.0	
4	2	1	1	7.60	14	6.0	

```
In [ ]: price_features.to_csv('price_features.csv')
```

```
In [91]: df_1 = pd.read_csv('wheelseye_data.csv')
pd.set_option('display.max_columns',20)
print(df.shape)
```

```
(10000, 18)
```

```
In [92]: df_1.head()
```

```
Out[92]:
```

	order_id	customer_id	city_from	city_to	route_distance_km	vehicle_type	material_weight	tru
0	ORDI000	CVS9726	Chennai	Delhi	1144	Open Truck	1.61 Tons	
1	ORDI001	CVS7464	Chennai	Pune	1912	Open Truck	1.98 Tons	
2	ORDI002	CVS4668	Kanpur	Indore	1145	Open Truck	7.28 Tons	
3	ORDI003	CVS2876	Indore	Chennai	1247	Open Truck	8.69 Tons	
4	ORDI004	CVS4137	Indore	Kanpur	569	Small Pickup Truck	7.6 Tons	

```
In [94]: price_features.columns
```

```
Out[94]: Index(['city_from', 'city_to', 'vehicle_type', 'material_weight',
               'truck_length_ft', 'truck_height_ft', 'business_category'],
              dtype='object')
```

```
In [95]: df_v1 = df_1[['city_from', 'city_to', 'vehicle_type', 'material_weight', 'tr
```

```
In [96]: df_v1.head()
```

```
Out[96]:
```

	city_from	city_to	vehicle_type	material_weight	truck_length_ft	truck_height_ft	business_cate
0	Chennai	Delhi	Open Truck	1.61 Tons	8	6.0	Transpo
1	Chennai	Pune	Open Truck	1.98 Tons	20	8.0	Transpo
2	Kanpur	Indore	Open Truck	7.28 Tons	7	6.0	Manufact
3	Indore	Chennai	Open Truck	8.69 Tons	19	8.0	Truck O
4	Indore	Kanpur	Small Pickup Truck	7.6 Tons	14	6.0	Transpo

```
In [97]: price_features.to_csv('encoded_features.csv')  
df_v1.to_csv('cc_features.csv')
```

Thank You!

```
In [99]: # More Advance Working Sonn..!
```

```
In [100... # Notebook Project By : PRASAD JADHAV (ML-ENG)
```

```
In [ ]: # LinkedIn: linkedin.com/in/prasadmjadhav2 Github: github.com/prasadmjadhav2
```